

train_fashion

May 14, 2019

1 Mask R-CNN - Train on Fashion Dataset

This notebook shows how to train Mask R-CNN on Modanet dataset. It is better to have a GPU, though, because the network backbone is a Resnet101, which would be too slow to train on a CPU. On a GPU, you can start to get okay-ish results in a few minutes, and good results in less than an hour.

```
In [1]: import os
import sys
import random
import math
import re
import time
import numpy as np
import cv2
import matplotlib
import matplotlib.pyplot as plt

# Root directory of the project
ROOT_DIR = os.path.abspath("../..")

# Import Mask RCNN
sys.path.append(ROOT_DIR) # To find local version of the library
from mrcnn.config import Config
from mrcnn import utils
import mrcnn.model as modellib
from mrcnn import visualize
from mrcnn.model import log

from pycocotools.coco import COCO

%matplotlib inline

# Directory to save logs and trained model
MODEL_DIR = os.path.join(ROOT_DIR, "logs")

# Local path to trained weights file
```

```

COCO_MODEL_PATH = os.path.join(ROOT_DIR, "mask_rcnn_coco.h5")
# Download COCO trained weights from Releases if needed
if not os.path.exists(COCO_MODEL_PATH):
    utils.download_trained_weights(COCO_MODEL_PATH)

```

Using TensorFlow backend.

1.1 Configurations

```

In [2]: class FashionConfig(Config):
        """Configuration for training on the toy shapes dataset.
        Derives from the base Config class and overrides values specific
        to the toy shapes dataset.
        """

        # Give the configuration a recognizable name
        NAME = "fashion"

        # Train on 1 GPU and 8 images per GPU. We can put multiple images on each
        # GPU because the images are small. Batch size is 8 (GPUs * images/GPU).
        GPU_COUNT = 1
        IMAGES_PER_GPU = 1

        # Number of classes (including background)
        NUM_CLASSES = 1 + 13 # background + 3 shapes

        # Use small images for faster training. Set the limits of the small side
        # the large side, and that determines the image shape.
        IMAGE_MIN_DIM = 128
        IMAGE_MAX_DIM = 128

        # Use smaller anchors because our image and objects are small
        RPN_ANCHOR_SCALES = (8, 16, 32, 64, 128) # anchor side in pixels

        # Reduce training ROIs per image because the images are small and have
        # few objects. Aim to allow ROI sampling to pick 33% positive ROIs.
        TRAIN_ROIS_PER_IMAGE = 32

        # Use a small epoch since the data is simple
        STEPS_PER_EPOCH = 100

        # use small validation steps since the epoch is small
        VALIDATION_STEPS = 5

        config = FashionConfig()
        config.display()

```

Configurations:

BACKBONE	resnet101
BACKBONE_STRIDES	[4, 8, 16, 32, 64]
BATCH_SIZE	1
BBOX_STD_DEV	[0.1 0.1 0.2 0.2]
COMPUTE_BACKBONE_SHAPE	None
DETECTION_MAX_INSTANCES	100
DETECTION_MIN_CONFIDENCE	0.7
DETECTION_NMS_THRESHOLD	0.3
FPN_CLASSIF_FC_LAYERS_SIZE	1024
GPU_COUNT	1
GRADIENT_CLIP_NORM	5.0
IMAGES_PER_GPU	1
IMAGE_CHANNEL_COUNT	3
IMAGE_MAX_DIM	128
IMAGE_META_SIZE	26
IMAGE_MIN_DIM	128
IMAGE_MIN_SCALE	0
IMAGE_RESIZE_MODE	square
IMAGE_SHAPE	[128 128 3]
LEARNING_MOMENTUM	0.9
LEARNING_RATE	0.001
LOSS_WEIGHTS	{'rpn_class_loss': 1.0, 'rpn_bbox_loss': 1.0, 'mrcnn_class_loss':
MASK_POOL_SIZE	14
MASK_SHAPE	[28, 28]
MAX_GT_INSTANCES	100
MEAN_PIXEL	[123.7 116.8 103.9]
MINI_MASK_SHAPE	(56, 56)
NAME	fashion
NUM_CLASSES	14
POOL_SIZE	7
POST_NMS_ROIS_INFERENCE	1000
POST_NMS_ROIS_TRAINING	2000
PRE_NMS_LIMIT	6000
ROI_POSITIVE_RATIO	0.33
RPN_ANCHOR_RATIOS	[0.5, 1, 2]
RPN_ANCHOR_SCALES	(8, 16, 32, 64, 128)
RPN_ANCHOR_STRIDE	1
RPN_BBOX_STD_DEV	[0.1 0.1 0.2 0.2]
RPN_NMS_THRESHOLD	0.7
RPN_TRAIN_ANCHORS_PER_IMAGE	256
STEPS_PER_EPOCH	100
TOP_DOWN_PYRAMID_SIZE	256
TRAIN_BN	False
TRAIN_ROIS_PER_IMAGE	32
USE_MINI_MASK	True
USE_RPN_ROIS	True
VALIDATION_STEPS	5
WEIGHT_DECAY	0.0001

1.2 Notebook Preferences

```
In [3]: def get_ax(rows=1, cols=1, size=8):
        """Return a Matplotlib Axes array to be used in
        all visualizations in the notebook. Provide a
        central point to control graph sizes.

        Change the default size attribute to control the size
        of rendered images
        """
        _, ax = plt.subplots(rows, cols, figsize=(size*cols, size*rows))
        return ax
```

1.3 Dataset

Extend the Dataset class and add a method to load the shapes dataset, `load_shapes()`, and override the following methods:

- `load_image()`
- `load_mask()`
- `image_reference()`

```
In [4]: import io
        import lmdb
        import sqlite3
        import pandas as pd
        import json
        from PIL import Image
        from IPython.display import display

In [7]: class PhotoData(object):
        def __init__(self, path):
            self.env = lmdb.open(
                path, map_size=2**36, readonly=True, lock=False
            )

        def __iter__(self):
            with self.env.begin() as t:
                with t.cursor() as c:
                    for key, value in c:
                        yield key, value

        def __getitem__(self, index):
            key = str(index).encode('ascii')
            with self.env.begin() as t:
```

```

        data = t.get(key)
    if not data:
        return None
    with io.BytesIO(data) as f:
        image = Image.open(f)
        image.load()
        return image

    def __len__(self):
        return self.env.stat()['entries']

photo_data = PhotoData(r'..\..\..\photos.lmdb')
print(len(photo_data))

```

1097474

In [8]: from pycocotools import mask as maskUtils

```

class FashionDataset(utils.Dataset):

    def load_fashion(self, count=5, start=0):
        json_file = r'..\..\..\modanet2018_instances_train.json'
        d = json.load(open(json_file))
        coco=COCO(json_file)
        class_ids = sorted(coco.getCatIds())
        for id in class_ids:
            self.add_class("fashion", id, "")

        image_ids = []
        for c in range(count):
            if c < 5:
                print(d['images'][c+start]['id'])
                image_ids.append(d['images'][c+start]['id'])

        # Add images
        for i in image_ids:
            annIds = coco.getAnnIds(imgIds=i, catIds=class_ids, iscrowd=None)
            anns = coco.loadAnns(annIds)
            width = coco.imgs[i]["width"]
            height = coco.imgs[i]["height"]
            self.add_image("fashion", image_id=i, path=None,
                           width=width, height=height, annotations=anns)

    def load_image(self, image_id):
        imgId = self.image_info[image_id]['id']
        image = photo_data[imgId]

```

```

out = np.array(image.getdata()).astype(np.int32).reshape((image.size[1], image.size[0]))
return out

def image_reference(self, image_id):
    """Return the shapes data of the image."""
    pass

def load_mask(self, image_id):
    """Load instance masks for the given image.

    Different datasets use different ways to store masks. This
    function converts the different mask format to one format
    in the form of a bitmap [height, width, instances].

    Returns:
    masks: A bool array of shape [height, width, instance count] with
           one mask per instance.
    class_ids: a 1D array of class IDs of the instance masks.
    """
    # If not a COCO image, delegate to parent class.
    image_info = self.image_info[image_id]

    instance_masks = []
    class_ids = []
    annotations = self.image_info[image_id]["annotations"]
    # Build mask of shape [height, width, instance_count] and list
    # of class IDs that correspond to each channel of the mask.
    for annotation in annotations:
        class_id = annotation['category_id']
        if class_id:
            m = self.annToMask(annotation, image_info["height"],
                               image_info["width"])
            # Some objects are so small that they're less than 1 pixel area
            # and end up rounded out. Skip those objects.
            if m.max() < 1:
                continue
            # Is it a crowd? If so, use a negative class ID.
            if annotation['iscrowd']:
                # Use negative class ID for crowds
                class_id *= -1
                # For crowd masks, annToMask() sometimes returns a mask
                # smaller than the given dimensions. If so, resize it.
                if m.shape[0] != image_info["height"] or m.shape[1] != image_info["width"]:
                    m = np.ones([image_info["height"], image_info["width"]], dtype=bool)
            instance_masks.append(m)
            class_ids.append(class_id)

```

```

        # Pack instance masks into an array
        if class_ids:
            mask = np.stack(instance_masks, axis=2).astype(np.bool)
            class_ids = np.array(class_ids, dtype=np.int32)
            return mask, class_ids
        else:
            # Call super class to return an empty mask
            return super(CocoDataset, self).load_mask(image_id)

def annToRLE(self, ann, height, width):
    """
    Convert annotation which can be polygons, uncompressed RLE to RLE.
    :return: binary mask (numpy 2D array)
    """
    segm = ann['segmentation']
    if isinstance(segm, list):
        # polygon -- a single object might consist of multiple parts
        # we merge all parts into one mask rle code
        rles = maskUtils.frPyObjects(segm, height, width)
        rle = maskUtils.merge(rles)
    elif isinstance(segm['counts'], list):
        # uncompressed RLE
        rle = maskUtils.frPyObjects(segm, height, width)
    else:
        # rle
        rle = ann['segmentation']
    return rle

def annToMask(self, ann, height, width):
    """
    Convert annotation which can be polygons, uncompressed RLE, or RLE to binary mask
    :return: binary mask (numpy 2D array)
    """
    rle = self.annToRLE(ann, height, width)
    m = maskUtils.decode(rle)
    return m

```

```

In [25]: # Training dataset
train_count = 1
val_count = 1
dataset_train = FashionDataset()
dataset_train.load_fashion(train_count)
dataset_train.prepare()

# Validation dataset - overfit 1 image
dataset_val = FashionDataset()

```

```

dataset_val.load_fashion(val_count)
dataset_val.prepare()

loading annotations into memory...
Done (t=3.52s)
creating index...
index created!
736791
loading annotations into memory...
Done (t=3.61s)
creating index...
index created!
736791

```

```

In [26]: # Load and display random samples
image_ids = np.random.choice(dataset_train.image_ids, 1)
for image_id in image_ids:
    image = dataset_train.load_image(image_id)
    mask, class_ids = dataset_train.load_mask(image_id)
    visualize.display_top_masks(image, mask, class_ids, dataset_train.class_names)

```



1.4 More data samples

```

In [27]: %%javascript
IPython.OutputArea.prototype._should_scroll = function(lines) {
    return false;
}

```

<IPython.core.display.Javascript object>

```

In [28]: sample_data = FashionDataset()
sample_data.load_fashion(10, start=10)
sample_data.prepare()
for image_id in sample_data.image_ids:
    image = sample_data.load_image(image_id)
    mask, class_ids = sample_data.load_mask(image_id)
    visualize.display_top_masks(image, mask, class_ids, sample_data.class_names)

```


loading annotations into memory...

Done (t=3.69s)

creating index...

index created!

764947

765216

765262

765317

765326

H x W=600x400



H x W=600x400



H x W=600x400



H x W=600x400







1.5 Create Model

```
In [29]: # Create model in training mode
        model = modellib.MaskRCNN(mode="training", config=config,
                                   model_dir=MODEL_DIR)

In [30]: # Which weights to start with?
        init_with = "coco" # imagenet, coco, or last

        if init_with == "imagenet":
            model.load_weights(model.get_imagenet_weights(), by_name=True)
        elif init_with == "coco":
            # Load weights trained on MS COCO, but skip layers that
            # are different due to the different number of classes
            # See README for instructions to download the COCO weights
            model.load_weights(COCO_MODEL_PATH, by_name=True,
                              exclude=["mrcnn_class_logits", "mrcnn_bbox_fc",
                                       "mrcnn_bbox", "mrcnn_mask"])
        elif init_with == "last":
            # Load the last model you trained and continue training
            model.load_weights(model.find_last(), by_name=True)
```

1.6 Training

Train in two stages: 1. Only the heads. Here we're freezing all the backbone layers and training only the randomly initialized layers (i.e. the ones that we didn't use pre-trained weights from MS COCO). To train only the head layers, pass `layers='heads'` to the `train()` function.

2. Fine-tune all layers. For this simple example it's not necessary, but we're including it to show the process. Simply pass `layers='all'` to train all layers.

```
In [31]: # Train the head branches
        # Passing layers="heads" freezes all layers except the head
        # layers. You can also pass a regular expression to select
        # which layers to train by name pattern.
```

```

model.train(dataset_train, dataset_val,
            learning_rate=config.LEARNING_RATE,
            epochs=1,
            layers='heads')

```

Starting at epoch 0. LR=0.001

Checkpoint Path: C:\Training\DeepLearningGit\virtualFitting\mask_rcnn\logs\fashion20190514T0059\
Selecting layers to train

```

fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)

```

In model: rpn_model

```

    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
    rpn_bbox_pred        (Conv2D)
mrcnn_mask_conv1      (TimeDistributed)
mrcnn_mask_bn1        (TimeDistributed)
mrcnn_mask_conv2      (TimeDistributed)
mrcnn_mask_bn2        (TimeDistributed)
mrcnn_class_conv1     (TimeDistributed)
mrcnn_class_bn1       (TimeDistributed)
mrcnn_mask_conv3      (TimeDistributed)
mrcnn_mask_bn3        (TimeDistributed)
mrcnn_class_conv2     (TimeDistributed)
mrcnn_class_bn2       (TimeDistributed)
mrcnn_mask_conv4      (TimeDistributed)
mrcnn_mask_bn4        (TimeDistributed)
mrcnn_bbox_fc         (TimeDistributed)
mrcnn_mask_deconv     (TimeDistributed)
mrcnn_class_logits    (TimeDistributed)
mrcnn_mask            (TimeDistributed)

```

C:\anaconda3\envs\mask_rcnn\lib\site-packages\tensorflow\python\ops\gradients_impl.py:110: UserWarning
"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

Epoch 1/1

100/100 [=====] - 232s 2s/step - loss: 1.3216 - rpn_class_loss: 0.0565

```

In [32]: # Fine tune all layers
         # Passing layers="all" trains all layers. You can also

```

```

# pass a regular expression to select which layers to
# train by name pattern.
model.train(dataset_train, dataset_val,
            learning_rate=config.LEARNING_RATE / 10,
            epochs=1,
            layers="all")

```

Starting at epoch 1. LR=0.0001

Checkpoint Path: C:\Training\DeepLearningGit\virtualFitting\mask_rcnn\logs\fashion20190514T0059\

Selecting layers to train

conv1	(Conv2D)
bn_conv1	(BatchNorm)
res2a_branch2a	(Conv2D)
bn2a_branch2a	(BatchNorm)
res2a_branch2b	(Conv2D)
bn2a_branch2b	(BatchNorm)
res2a_branch2c	(Conv2D)
res2a_branch1	(Conv2D)
bn2a_branch2c	(BatchNorm)
bn2a_branch1	(BatchNorm)
res2b_branch2a	(Conv2D)
bn2b_branch2a	(BatchNorm)
res2b_branch2b	(Conv2D)
bn2b_branch2b	(BatchNorm)
res2b_branch2c	(Conv2D)
bn2b_branch2c	(BatchNorm)
res2c_branch2a	(Conv2D)
bn2c_branch2a	(BatchNorm)
res2c_branch2b	(Conv2D)
bn2c_branch2b	(BatchNorm)
res2c_branch2c	(Conv2D)
bn2c_branch2c	(BatchNorm)
res3a_branch2a	(Conv2D)
bn3a_branch2a	(BatchNorm)
res3a_branch2b	(Conv2D)
bn3a_branch2b	(BatchNorm)
res3a_branch2c	(Conv2D)
res3a_branch1	(Conv2D)
bn3a_branch2c	(BatchNorm)
bn3a_branch1	(BatchNorm)
res3b_branch2a	(Conv2D)
bn3b_branch2a	(BatchNorm)
res3b_branch2b	(Conv2D)
bn3b_branch2b	(BatchNorm)
res3b_branch2c	(Conv2D)
bn3b_branch2c	(BatchNorm)

res3c_branch2a	(Conv2D)
bn3c_branch2a	(BatchNorm)
res3c_branch2b	(Conv2D)
bn3c_branch2b	(BatchNorm)
res3c_branch2c	(Conv2D)
bn3c_branch2c	(BatchNorm)
res3d_branch2a	(Conv2D)
bn3d_branch2a	(BatchNorm)
res3d_branch2b	(Conv2D)
bn3d_branch2b	(BatchNorm)
res3d_branch2c	(Conv2D)
bn3d_branch2c	(BatchNorm)
res4a_branch2a	(Conv2D)
bn4a_branch2a	(BatchNorm)
res4a_branch2b	(Conv2D)
bn4a_branch2b	(BatchNorm)
res4a_branch2c	(Conv2D)
res4a_branch1	(Conv2D)
bn4a_branch2c	(BatchNorm)
bn4a_branch1	(BatchNorm)
res4b_branch2a	(Conv2D)
bn4b_branch2a	(BatchNorm)
res4b_branch2b	(Conv2D)
bn4b_branch2b	(BatchNorm)
res4b_branch2c	(Conv2D)
bn4b_branch2c	(BatchNorm)
res4c_branch2a	(Conv2D)
bn4c_branch2a	(BatchNorm)
res4c_branch2b	(Conv2D)
bn4c_branch2b	(BatchNorm)
res4c_branch2c	(Conv2D)
bn4c_branch2c	(BatchNorm)
res4d_branch2a	(Conv2D)
bn4d_branch2a	(BatchNorm)
res4d_branch2b	(Conv2D)
bn4d_branch2b	(BatchNorm)
res4d_branch2c	(Conv2D)
bn4d_branch2c	(BatchNorm)
res4e_branch2a	(Conv2D)
bn4e_branch2a	(BatchNorm)
res4e_branch2b	(Conv2D)
bn4e_branch2b	(BatchNorm)
res4e_branch2c	(Conv2D)
bn4e_branch2c	(BatchNorm)
res4f_branch2a	(Conv2D)
bn4f_branch2a	(BatchNorm)
res4f_branch2b	(Conv2D)
bn4f_branch2b	(BatchNorm)

res4f_branch2c	(Conv2D)
bn4f_branch2c	(BatchNorm)
res4g_branch2a	(Conv2D)
bn4g_branch2a	(BatchNorm)
res4g_branch2b	(Conv2D)
bn4g_branch2b	(BatchNorm)
res4g_branch2c	(Conv2D)
bn4g_branch2c	(BatchNorm)
res4h_branch2a	(Conv2D)
bn4h_branch2a	(BatchNorm)
res4h_branch2b	(Conv2D)
bn4h_branch2b	(BatchNorm)
res4h_branch2c	(Conv2D)
bn4h_branch2c	(BatchNorm)
res4i_branch2a	(Conv2D)
bn4i_branch2a	(BatchNorm)
res4i_branch2b	(Conv2D)
bn4i_branch2b	(BatchNorm)
res4i_branch2c	(Conv2D)
bn4i_branch2c	(BatchNorm)
res4j_branch2a	(Conv2D)
bn4j_branch2a	(BatchNorm)
res4j_branch2b	(Conv2D)
bn4j_branch2b	(BatchNorm)
res4j_branch2c	(Conv2D)
bn4j_branch2c	(BatchNorm)
res4k_branch2a	(Conv2D)
bn4k_branch2a	(BatchNorm)
res4k_branch2b	(Conv2D)
bn4k_branch2b	(BatchNorm)
res4k_branch2c	(Conv2D)
bn4k_branch2c	(BatchNorm)
res4l_branch2a	(Conv2D)
bn4l_branch2a	(BatchNorm)
res4l_branch2b	(Conv2D)
bn4l_branch2b	(BatchNorm)
res4l_branch2c	(Conv2D)
bn4l_branch2c	(BatchNorm)
res4m_branch2a	(Conv2D)
bn4m_branch2a	(BatchNorm)
res4m_branch2b	(Conv2D)
bn4m_branch2b	(BatchNorm)
res4m_branch2c	(Conv2D)
bn4m_branch2c	(BatchNorm)
res4n_branch2a	(Conv2D)
bn4n_branch2a	(BatchNorm)
res4n_branch2b	(Conv2D)
bn4n_branch2b	(BatchNorm)

res4n_branch2c	(Conv2D)
bn4n_branch2c	(BatchNorm)
res4o_branch2a	(Conv2D)
bn4o_branch2a	(BatchNorm)
res4o_branch2b	(Conv2D)
bn4o_branch2b	(BatchNorm)
res4o_branch2c	(Conv2D)
bn4o_branch2c	(BatchNorm)
res4p_branch2a	(Conv2D)
bn4p_branch2a	(BatchNorm)
res4p_branch2b	(Conv2D)
bn4p_branch2b	(BatchNorm)
res4p_branch2c	(Conv2D)
bn4p_branch2c	(BatchNorm)
res4q_branch2a	(Conv2D)
bn4q_branch2a	(BatchNorm)
res4q_branch2b	(Conv2D)
bn4q_branch2b	(BatchNorm)
res4q_branch2c	(Conv2D)
bn4q_branch2c	(BatchNorm)
res4r_branch2a	(Conv2D)
bn4r_branch2a	(BatchNorm)
res4r_branch2b	(Conv2D)
bn4r_branch2b	(BatchNorm)
res4r_branch2c	(Conv2D)
bn4r_branch2c	(BatchNorm)
res4s_branch2a	(Conv2D)
bn4s_branch2a	(BatchNorm)
res4s_branch2b	(Conv2D)
bn4s_branch2b	(BatchNorm)
res4s_branch2c	(Conv2D)
bn4s_branch2c	(BatchNorm)
res4t_branch2a	(Conv2D)
bn4t_branch2a	(BatchNorm)
res4t_branch2b	(Conv2D)
bn4t_branch2b	(BatchNorm)
res4t_branch2c	(Conv2D)
bn4t_branch2c	(BatchNorm)
res4u_branch2a	(Conv2D)
bn4u_branch2a	(BatchNorm)
res4u_branch2b	(Conv2D)
bn4u_branch2b	(BatchNorm)
res4u_branch2c	(Conv2D)
bn4u_branch2c	(BatchNorm)
res4v_branch2a	(Conv2D)
bn4v_branch2a	(BatchNorm)
res4v_branch2b	(Conv2D)
bn4v_branch2b	(BatchNorm)

res4v_branch2c	(Conv2D)
bn4v_branch2c	(BatchNorm)
res4w_branch2a	(Conv2D)
bn4w_branch2a	(BatchNorm)
res4w_branch2b	(Conv2D)
bn4w_branch2b	(BatchNorm)
res4w_branch2c	(Conv2D)
bn4w_branch2c	(BatchNorm)
res5a_branch2a	(Conv2D)
bn5a_branch2a	(BatchNorm)
res5a_branch2b	(Conv2D)
bn5a_branch2b	(BatchNorm)
res5a_branch2c	(Conv2D)
res5a_branch1	(Conv2D)
bn5a_branch2c	(BatchNorm)
bn5a_branch1	(BatchNorm)
res5b_branch2a	(Conv2D)
bn5b_branch2a	(BatchNorm)
res5b_branch2b	(Conv2D)
bn5b_branch2b	(BatchNorm)
res5b_branch2c	(Conv2D)
bn5b_branch2c	(BatchNorm)
res5c_branch2a	(Conv2D)
bn5c_branch2a	(BatchNorm)
res5c_branch2b	(Conv2D)
bn5c_branch2b	(BatchNorm)
res5c_branch2c	(Conv2D)
bn5c_branch2c	(BatchNorm)
fpn_c5p5	(Conv2D)
fpn_c4p4	(Conv2D)
fpn_c3p3	(Conv2D)
fpn_c2p2	(Conv2D)
fpn_p5	(Conv2D)
fpn_p2	(Conv2D)
fpn_p3	(Conv2D)
fpn_p4	(Conv2D)
In model: rpn_model	
rpn_conv_shared	(Conv2D)
rpn_class_raw	(Conv2D)
rpn_bbox_pred	(Conv2D)
mrcnn_mask_conv1	(TimeDistributed)
mrcnn_mask_bn1	(TimeDistributed)
mrcnn_mask_conv2	(TimeDistributed)
mrcnn_mask_bn2	(TimeDistributed)
mrcnn_class_conv1	(TimeDistributed)
mrcnn_class_bn1	(TimeDistributed)
mrcnn_mask_conv3	(TimeDistributed)
mrcnn_mask_bn3	(TimeDistributed)

```

mrcnn_class_conv2      (TimeDistributed)
mrcnn_class_bn2        (TimeDistributed)
mrcnn_mask_conv4       (TimeDistributed)
mrcnn_mask_bn4         (TimeDistributed)
mrcnn_bbox_fc          (TimeDistributed)
mrcnn_mask_deconv      (TimeDistributed)
mrcnn_class_logits     (TimeDistributed)
mrcnn_mask              (TimeDistributed)

```

```

C:\anaconda3\envs\mask_rcnn\lib\site-packages\tensorflow\python\ops\gradients_impl.py:110: UserWarning:
  "Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

```

```

In [33]: # Save weights
         # Typically not needed because callbacks save after every epoch
         # Uncomment to save manually
         # model_path = os.path.join(MODEL_DIR, "mask_rcnn_fashion.h5")
         # model.keras_model.save_weights(model_path)

```

1.7 Detection

```

In [34]: class InferenceConfig(FashionConfig):

```

```

        GPU_COUNT = 1
        IMAGES_PER_GPU = 1

```

```

inference_config = InferenceConfig()

```

```

# Recreate the model in inference mode
model = modellib.MaskRCNN(mode="inference",
                           config=inference_config,
                           model_dir=MODEL_DIR)

```

```

# Get path to saved weights
# Either set a specific path or find last trained weights
# model_path = os.path.join(ROOT_DIR, ".h5 file name here")
model_path = model.find_last()

```

```

# Load trained weights
print("Loading weights from ", model_path)
model.load_weights(model_path, by_name=True)

```

```

Loading weights from C:\Training\DeepLearningGit\virtualFitting\mask_rcnn\logs\fashion20190514T
Re-starting from epoch 1

```

```

In [35]: # Test on a random image
         # image_id = random.choice(dataset_val.image_ids)
         # image_ids = dataset_val.image_ids

```

```


image_ids = [0]
for image_id in image_ids:
    original_image, image_meta, gt_class_id, gt_bbox, gt_mask = \
        modellib.load_image_gt(dataset_val, inference_config,
                                image_id, use_mini_mask=False)

    log("original_image", original_image)
    log("image_meta", image_meta)
    log("gt_class_id", gt_class_id)
    log("gt_bbox", gt_bbox)
    log("gt_mask", gt_mask)

    visualize.display_instances(original_image, gt_bbox, gt_mask, gt_class_id,
                               dataset_train.class_names, figsize=(8, 8))

```

original_image	shape: (128, 128, 3)	min:	0.00000	max:	255.00000	int32
image_meta	shape: (26,)	min:	0.00000	max:	600.00000	float64
gt_class_id	shape: (5,)	min:	2.00000	max:	11.00000	int32
gt_bbox	shape: (5, 4)	min:	28.00000	max:	124.00000	int32
gt_mask	shape: (128, 128, 5)	min:	0.00000	max:	1.00000	bool



```
In [36]: results = model.detect([original_image], verbose=1)
```

```

r = results[0]
visualize.display_instances(original_image, r['rois'], r['masks'], r['class_ids'],
                           dataset_val.class_names, r['scores'], ax=get_ax())

```

Processing 1 images

image	shape: (128, 128, 3)	min: 0.00000	max: 255.00000	int32
molded_images	shape: (1, 128, 128, 3)	min: -123.70000	max: 151.10000	float64
image metas	shape: (1, 26)	min: 0.00000	max: 128.00000	int32
anchors	shape: (1, 4092, 4)	min: -0.71267	max: 1.20874	float32



1.8 Evaluation

```
In [37]: # Compute VOC-Style mAP @ IoU=0.5
# Running on 10 images. Increase for better accuracy.
image_ids = np.random.choice(dataset_val.image_ids, 10)
APs = []
for image_id in image_ids:
    # Load image and ground truth data
    image, image_meta, gt_class_id, gt_bbox, gt_mask = \
        modellib.load_image_gt(dataset_val, inference_config,
                                image_id, use_mini_mask=False)
    molded_images = np.expand_dims(modellib.mold_image(image, inference_config), 0)
    # Run object detection
```

```

results = model.detect([image], verbose=0)
r = results[0]
# Compute AP
AP, precisions, recalls, overlaps =\
    utils.compute_ap(gt_bbox, gt_class_id, gt_mask,
                     r["rois"], r["class_ids"], r["scores"], r['masks'])
APs.append(AP)

print("mAP: ", np.mean(APs))

mAP: 0.5800000101327896

```

1.9 Applications once get the mask

Here is one simple example: if we know the pixel region of the skirt and top, we can change the textures or perform any other operations in those regions.

```

In [152]: # skirt is 11
print(r['masks'].shape)
print(r['class_ids'])
skirt = np.where(r['class_ids']==11)
skirt_idx = skirt[0][0]
mask_skirt = r['masks'][:, :, skirt_idx]
mask_top = r['masks'][:, :, np.where(r['class_ids']==9)[0][0]]

(128, 128, 18)
[ 9 11  4  4  4  4  4  4  4  2  4  4  4  4  4  4  4  4]

```

```

In [160]: import matplotlib.pyplot as plt
          %matplotlib inline
          import pylab
          import numpy as np

          image = original_image
          out = np.array(image, np.int32)
          # mask out skirt
          other = out * (mask_skirt[:, :, None] == 0)
          skirt = out * (mask_skirt[:, :, None] != 0)
          change = (out-90) * (mask_skirt[:, :, None] != 0)
          mix = other + skirt

          plt.figure(1)
          plt.subplot(1,2,1).axis('off')
          plt.imshow(mask_skirt)
          plt.subplot(1,2,2).axis('off')
          plt.imshow(mix)

```

```

plt.figure(2)
plt.subplot(1,2,1).axis('off')
plt.imshow(skirt)
plt.subplot(1,2,2).axis('off')
plt.imshow(other+change)

out = other + change
other = out * (mask_top[:, :, None] == 0)
top = out * (mask_top[:, :, None] != 0)
change = (out+100) * (mask_top[:, :, None] != 0)
mix = other + top

plt.figure(3)
plt.subplot(1,2,1).axis('off')
plt.imshow(mask_top)
plt.subplot(1,2,2).axis('off')
plt.imshow(mix)

plt.figure(4)
plt.subplot(1,2,1).axis('off')
plt.imshow(top)
plt.subplot(1,2,2).axis('off')
plt.imshow(other+change)

```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] f
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] f
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] f

Out[160]: <matplotlib.image.AxesImage at 0x1ff54a130b8>



